

Approximation: Greedy and Local Search

Zixuan Fan

Technische Universität München

July 2024

Agenda

1. Introduction
2. Scheduling Problems
 - ▶ Scheduling with deadlines on a Single machine
 - ▶ Scheduling on Multiple Machines
3. Graph Problems
 - ▶ K-center Problem
 - ▶ Travelling Salesman Problem
4. Conclusion

Purpose of Approximation

- ▶ Some problems not efficiently solvable: NP-hardness

Purpose of Approximation

- ▶ Some problems not efficiently solvable: NP-hardness
- ▶ Can we find a flawed solution?

Purpose of Approximation

- ▶ Some problems not efficiently solvable: NP-hardness
- ▶ Can we find a flawed solution?
- ▶ How flawed/good is this solution?

Purpose of Approximation

- ▶ Some problems not efficiently solvable: NP-hardness
- ▶ Can we find a flawed solution?
- ▶ How flawed/good is this solution?
- ▶ What is the limit of this flawed solution?

Quick Recap: Approximation Ratio

Opt^* : the optimal solution

Opt : the suboptimal solution we compute

For minimization problems, we have

$$\alpha = \frac{|Opt|}{|Opt^*|} > 1$$

And for maximization problems

$$\alpha = \frac{|Opt|}{|Opt^*|} < 1$$

Techniques in Approximation

1. Randomization: MAXSAT, MAXCUT

Techniques in Approximation

1. Randomization: MAXSAT, MAXCUT
2. Dynamic Programming: Knapsack, Bin Packing

Techniques in Approximation

1. Randomization: MAXSAT, MAXCUT
2. Dynamic Programming: Knapsack, Bin Packing
3. Linear & Integer Programming: Primal-Dual, Semidefinite Program

Techniques in Approximation

1. Randomization: MAXSAT, MAXCUT
2. Dynamic Programming: Knapsack, Bin Packing
3. Linear & Integer Programming: Primal-Dual, Semidefinite Program
4. Greedy & Local Search

Greedy and Local Search

- ▶ Both strategies attempt to make the best decision

Greedy and Local Search

- ▶ Both strategies attempt to make the best decision
- ▶ Greedy algorithms forms a solution step by step

Greedy and Local Search

- ▶ Both strategies attempt to make the best decision
- ▶ Greedy algorithms forms a solution step by step
- ▶ Local Search starts search from an arbitrary solution

Scheduling with deadlines on a Single machine

Problem Statement: Given n jobs to be processed on a single machine. How can we schedule them such that they will finish as early as possible.

- ▶ What if there is not deadline?
- ▶ How do we define earliness/lateness?

Scheduling with deadlines on a Single machine - Formal Definition

Suppose job j is finished at time C_j , the lateness is

$$L_j := C_j - d_j$$

The lateness of all jobs is

$$L_{max} = \max_{i \in [n]} L_i$$

Input: n jobs with release time r_j , processing time p_j , and deadline d_j .

Output: a schedule such that L_{max} is minimized.

Algorithm - Intuition

If the job j is finished before the deadline we don't get penalty for it, so ...

1. What if we always choose the job with the earliest deadlines

Algorithm - Intuition

If the job j is finished before the deadline we don't get penalty for it, so ...

1. What if we always choose the job with the earliest deadlines
2. But, we may have deal with negative values.

Algorithm - Intuition

If the job j is finished before the deadline we don't get penalty for it, so ...

1. What if we always choose the job with the earliest deadlines
2. But, we may have deal with negative values.
3. So all deadlines are negative (by assumption).

Algorithm - Analysis

We start with an observation. Let S be a set of jobs,

- ▶ $r(S) := \min_{j \in S} r_j$
- ▶ $p(S) := \sum_{j \in S} p_j$
- ▶ $d(S) := \max_{j \in S} d_j$

We claim that

$$L_{max}^* \geq r(S) + p(S) - d(S)$$

This is proven by considering the optimal schedule.

Algorithm - Analysis

$$L_{max}^* \geq r(S) + p(S) - d(S)$$

This leads directly to

$$L_{max}^* \geq r(j) + p(j) - d(j) \geq -d_j$$

where j is the job that leads to the maximal lateness.

Algorithm - Analysis

$$L_{max}^* \geq -d_j$$

Recall $L_{max} = C_j - d_j$, it suffices to show

$$L_{max}^* \geq C_j$$

Algorithm - Analysis

We start with an ideal scenario: all jobs are released at time $t = 0$.
We have

- ▶ $r(S) = 0$
- ▶ $p(S) = C_j$
- ▶ $d(S) < 0$, by assumption

From lemma it follows that

$$L_{max}^* \geq r(S) + p(S) - d(S) \geq p(S) = C_j$$

Algorithm - Analysis

We start with an ideal scenario: all jobs are released at time $t = 0$.
We have

- ▶ $r(S) = 0$
- ▶ $p(S) = C_j$
- ▶ $d(S) < 0$, by assumption

From lemma it follows that

$$L_{max}^* \geq r(S) + p(S) - d(S) \geq p(S) = C_j$$

How about the general cases?

Algorithm - Analysis

How about the general cases?

- ▶ We consider it in a way of calculus

Algorithm - Analysis

How about the general cases?

- ▶ We consider it in a way of calculus
- ▶ Find the time t such that $[t, C_j]$ has no idle time

Algorithm - Analysis

How about the general cases?

- ▶ We consider it in a way of calculus
- ▶ Find the time t such that $[t, C_j]$ has no idle time
- ▶ Denote jobs processed in this interval with S .

Algorithm - Analysis

How about the general cases?

- ▶ We consider it in a way of calculus
- ▶ Find the time t such that $[t, C_j]$ has no idle time
- ▶ Denote jobs processed in this interval with S .
- ▶ $r(S) = t$, $p(S) = C_j - t$

Algorithm - Analysis

How about the general cases?

- ▶ We consider it in a way of calculus
- ▶ Find the time t such that $[t, C_j]$ has no idle time
- ▶ Denote jobs processed in this interval with S .
- ▶ $r(S) = t$, $p(S) = C_j - t$
- ▶ Thus $C_j = r(S) + p(S) \leq L_{max}^*$

Algorithm - Analysis

Summarizing two results $L_{max}^* \geq -d_j$ and $L_{max}^* \geq C_j$, we obtain

$$L_{max} = C_j - d_j \leq 2L_{max}^*$$

Scheduling on Identical Parallel Machines

Problem Statement: Given n jobs to be processed on k machines, How can we schedule them such that they will finish as early as possible.

- ▶ What if there is not deadline?
- ▶ How do we define earliness/lateness?

Scheduling on Identical Parallel Machines - Formal Definition

More machines \implies more complexity

So the problem is **NP-hard** even if

we don't consider **release time** and **deadlines**

Scheduling on Identical Parallel Machines - Formal Definition

Input: n jobs with processing time p_j and k machines

Output: A schedule that minimizes

$$\max_{j \in [n]} C_j$$

where C_j is the completion time of job j .

Algorithm - Local Search

Recap: Local search

- ▶ starts with a valid solution

Algorithm - Local Search

Recap: Local search

- ▶ starts with a valid solution
- ▶ searches greedily until

Algorithm - Local Search

Recap: Local search

- ▶ starts with a valid solution
- ▶ searches greedily until
- ▶ no better strategy can be made

Algorithm - Local Search

Recap: Local search

- ▶ How do we find a valid solution?
- ▶ How to make greedy strategy?

Algorithm - Local Search

Recap: Local search

- ▶ How do we find a valid solution?
- ▶ How to make greedy strategy?

Algorithm - Local Search

Recap: Local search

- ▶ How do we find a valid solution?
- ▶ **Finish all jobs in one machine.**
- ▶ How to make greedy strategy?

Algorithm - Local Search

Recap: Local search

- ▶ How do we find a valid solution?
- ▶ Finish all jobs in one machine.
- ▶ How to make greedy strategy?
- ▶ Separate those jobs to other machines.

Local Search - Separation of Jobs

- ▶ Pick the last job j at the machine M with the latest running time
- ▶ Let t be the processing time of all other jobs on M
- ▶ Find another machine M' such that $t' < t$ is minimized
- ▶ Add j to M'

Local Search - Separation of Jobs

- ▶ Pick the last job j at the machine M with the latest running time
- ▶ Let t be the processing time of all other jobs on M
- ▶ Find another machine M' such that $t' < t$ is minimized
- ▶ Add j to M'

But, can we go into **LOOPS** ?

Local Search - Analysis

We show that a job is only assigned **once** in our local search algorithm.

Suppose that a job j on machine M is assigned for a **second** time.

Local Search - Analysis

We show that a job is only assigned **once** in our local search algorithm.

Suppose that a job j on machine M is assigned for a **second** time.

\implies machine M has the latest running time

Local Search - Analysis

We show that a job is only assigned **once** in our local search algorithm.

Suppose that a job j on machine M is assigned for a **second** time.

\implies machine M has the latest running time

\implies there is another machine M' such that $t' < t$

Local Search - Analysis

We show that a job is only assigned **once** in our local search algorithm.

Suppose that a job j on machine M is assigned for a **second** time.

\implies machine M has the latest running time

\implies there is another machine M' such that $t' < t$

BUT: by our strategy t has to be minimal when j was assigned at an earlier step.

Local Search - Analysis

We show that a job is only assigned **once** in our local search algorithm.

Suppose that a job j on machine M is assigned for a **second** time.

\implies machine M has the latest running time

\implies there is another machine M' such that $t' < t$

BUT: by our strategy t has to be minimal when j was assigned at an earlier step.

We denote completion time on M' at this step as t'_0 , it holds

$$t' \geq t'_0 \geq t \nmid$$

Local Search - Analysis

What about the approximation ratio?

Local Search - Analysis

What about the approximation ratio?

- ▶ Completion time on each machine should be close to mean:

$$\frac{\sum_{j \in [n]} p_j}{n}$$

- ▶ It's also the case for the optimal solution, and

$$|Opt^*| \geq \frac{\sum_{j \in [n]} p_j}{k}$$

- ▶ We show

$$|Opt| \leq \frac{2 \sum_{j \in [n]} p_j}{k} \leq 2|Opt^*|$$

Local Search - Analysis

We show

$$|Opt| \leq \frac{2 \sum_{j \in [n]} p_j}{k}$$

- ▶ Consider the machine M with the longest completion time
- ▶ Job j that completes last with processing time p_j
- ▶ All other jobs completes at t .

We have

$$|Opt| = t + p_j$$

and we may witness

$$t \leq \frac{\sum_{j \in [n]} p_j}{k}$$
$$p_j \leq \frac{\sum_{j \in [n]} p_j}{k}$$

Algorithm - Greedy and More

- ▶ There exists greedy algorithm that solves problem with $\alpha = 2, \frac{4}{3}$
- ▶ With rounding and DP, the problem can be approximated with $\alpha = 1 + \epsilon$ for any $\epsilon > 0$

K-Center Problem

Problem Statement: Given n points in a space, choose k centers from them such that the sum of distances between each point and its nearest neighbor is minimized.

K-Center Problem - Formal Definition

$d(p, q)$: distance between p, q

$d(p, S) := \min_{q \in S} d(p, q)$

Input: S, k

Output: $C \subseteq S$ s.t. $|C| = k$ and $\sum_{p \in S} d(p, S)$ is minimized

K-Center Problem - Formal Definition

$d(p, q)$: distance between p, q

$d(p, S) := \min_{q \in S} d(p, q)$

Input: S, k

Output: $C \subseteq S$ s.t. $|C| = k$ and $\sum_{p \in S} d(p, S)$ is minimized

What is the distance function?

Or what property does the distance function have?

Algorithm - Intuition

How do we search **greedily**?

If we have a temporary center $|C| < k$, how to build $|C'| = |C| + 1$?

Algorithm - Description

Find $v \in S \setminus C$ s.t. $d(v, C)$ is maximized.

$$C' = C \cup \{v\}$$

Algorithm - Description

Find $v \in S \setminus C$ s.t. $d(v, C)$ is maximized.

$$C' = C \cup \{v\}$$

What is the cost per iteration?

Algorithm - Description

Find $v \in S \setminus C$ s.t. $d(v, C)$ is maximized.

$$C' = C \cup \{v\}$$

What is the cost per iteration? $O(|S| \cdot |S \setminus V|) = O(n^2)$

Algorithm - Description

But, what is C at the beginning of the iteration?

Algorithm - Description

But, what is C at the beginning of the iteration?

1. The center of all points? \implies Needs another $O(n^2)$ cost
2. Any arbitrary point \implies Constant cost

Algorithm - Description

But, what is C at the beginning of the iteration?

1. The center of all points? \implies Needs another $O(n^2)$ cost
2. Any arbitrary point \implies Constant cost

Is the $O(n^2)$ cost worth it? We will discuss it later.

Analysis

Quadratic cost

Analysis

Quadratic cost

Approximation ratio?

Approximation ratio: Base

Suppose each optimal cluster contains exactly one center by our algorithm.

Approximation ratio: Base

Suppose each optimal cluster contains exactly one center by our algorithm.

If all points in this cluster is covered by the cluster of this new center

$$r \leq 2r^*$$

If they are not covered \implies a shorter radius is ok.

Approximation ratio: Base

Suppose each optimal cluster contains exactly one center by our algorithm.

If all points in this cluster is covered by the cluster of this new center

$$r \leq 2r^*$$

If they are not covered \implies a shorter radius is ok.

But guarantees $\alpha = \frac{r}{r^*} = 2$

Approximation ratio: General

What if one optimal cluster contains two or even more centers?

This is solve by our greediness. **HOW?**

Approximation ratio: Improvement?

Can we have a better approximation?

Sadly, NO

Traverlling Salesman Problem(TSP)

Problem Statement: Given a complete weighted graph of n vertices, find the round tour with the minimal cost.

TSP - Formal Definition

Input: $G = (V, d)$

Output: Path π s.t. $\sum_{i=1}^{n-1} w(\pi(i), \pi(i+1))$ is minimized

TSP - Formal Definition

Input: $G = (V, d)$

Output: Path π s.t. $\sum_{i=1}^{n-1} w(\pi(i), \pi(i+1))$ is minimized

What is the distance function?

Or what property does the distance function have?

TSP-Limit of Approximation

1. No existent approximation for TSP in general

TSP-Limit of Approximation

1. No existent approximation for TSP in general
2. If there was \implies Hamilton-Cycle solvable in P!

Suppose for each vertice $u, v, w \in V$, triangular inequality holds

$$d(u, v) + d(v, w) \geq d(u, w)$$

TSP-Limit of Approximation

1. No existent approximation for TSP in general
2. If there was \implies Hamilton-Cycle solvable in P!
3. But: metric distance can help us!

Suppose for each vertex $u, v, w \in V$, triangular inequality holds

$$d(u, v) + d(v, w) \geq d(u, w)$$

Algorithm - A bit help from Euler

- ▶ Minimum Spanning Tree(MST). The tree that connects all vertices with minimal edge cost.

Algorithm - A bit help from Euler

- ▶ Minimum Spanning Tree(MST). The tree that connects all vertices with minimal edge cost.
- ▶ Eulerian Graph: connected graph where all vertices have even degrees.

Algorithm - A bit help from Euler

- ▶ What if we make a copy of MST?

Algorithm - A bit help from Euler

- ▶ What if we make a copy of MST?
- ▶ This is make a Eulerian Graph \implies Not a round tour

Algorithm - A bit help from Euler

- ▶ What if we make a copy of MST?
- ▶ This is make a Eulerian Graph \implies Not a round tour
- ▶ What if we remove all but the first occurrences?

Algorithm - A bit help from Euler

- ▶ What if we make a copy of MST?
- ▶ This is make a Eulerian Graph \implies Not a round tour
- ▶ What if we remove all but the first occurrences?
- ▶ This is a round tour!

Algorithm - A bit help from Euler

For the size of the tour:

1. $w_{MST} \leq w_{Opt^*}$: Round tour contains trees!
2. $w_G \leq 2w_{MST}$: triangular inequality

$$\implies w_G \leq 2w_{Opt^*}$$

Algorithm - A bit help from Christofide

- ▶ We can do the same trick to all Eulerian Graphs

Algorithm - A bit help from Christofide

- ▶ We can do the same trick to all Eulerian Graphs
- ▶ As long as it connects all vertices

Algorithm - A bit help from Christofide

- ▶ We can do the same trick to all Eulerian Graphs
- ▶ As long as it connects all vertices
- ▶ How to make MST Eulerian?

Algorithm - A bit help from Christofide

- ▶ We can do the same trick to all Eulerian Graphs
- ▶ As long as it connects all vertices
- ▶ How to make MST Eulerian?
- ▶ Make all vertices have **even** degrees

Algorithm - A bit help from Christofide

HOW?

1. Connect all vertices with odd degrees

Algorithm - A bit help from Christofide

HOW?

1. Connect all vertices with odd degrees
2. How many vertices with odd degrees?

Algorithm - A bit help from Christofide

HOW?

1. Connect all vertices with odd degrees
2. How many vertices with odd degrees?
3. Finding a perfect matching

Algorithm - A bit help from Christofide

HOW?

1. Connect all vertices with odd degrees
2. How many vertices with odd degrees?
3. Finding a perfect matching
4. What is the cost? What is the size of the PM?

Algorithm - A bit help from Christofide

For the size of the tour:

1. $w_{MST} \leq w_{Opt^*}$: see 1st part
 2. $w_{PM} \leq \frac{1}{2}w_{Opt^*}$:
 3. $w_G \leq w_{MST} + w_{PM}$: traingular inequality
- $$\implies w_G \leq \frac{3}{2}w_{Opt^*}$$

Limit of Approximation(Again)

Even if it is possible to do the approximation for metric TSP, we cannot do better than $\alpha < \frac{220}{219}$ unless **P** = **NP**